

HW 06: k-NN Accelerator

Due: December 21st 23:59pm


Goal

In this assignment, you will implement k-nearest neighbors algorithm (k-NN) in C and Verilog on PicoRV32. The implementation details with PicoRV32 are the same as the previous assignment. But this time, the accelerator also acts as a master to the memory. You will need to read images out of the memory and process them.

Files Description

data_batch.bin: contains 10000 3073-byte rows of images in 10 classes from CIFAR-10 image dataset. Each image comes with a label from 0 to 9, indicating one of the ten classes. We only need the first 1000 images in this assignment. Here are the 10 classes in the dataset. For each class, we also list 10 random images for your reference:

airplane




automobile




bird



cat



deer




dog




frog




horse



ship



truck



golden.txt: The golden results (for $k=5$) to verify your results

imagegen.v: The tool to generate a bmp file of image #ID (0~999) for visualization, for example:

```
make image ID=0
```

Problem Description

k-NN: <http://cs231n.github.io/classification/#nn>

The k-NN algorithm could be summarized as follow:

1. Use the first 50 images as the test images, and the rest 950 images as the training images
2. For each test image, compute its distances to all the training images;
3. Find the k closest images (with the least distance);
4. Have these top k closest images vote on the label. Break tie by choosing the smaller class ID.

Memory Interface

The [native memory interface](#) of PicoRV32 is a simple valid-ready interface that can run one memory transfer at a time. Our accelerators and CPU could access the memory at the same time (similar to a **dual-port** memory). Furthermore, the Read-Cycle timing is just like the SRAM we used in **HW 03: FIFO**.

The 1000 images are stored in the memory from 0x00010000 in row-major order. An image row contains a label following by 3072 pixels ($32 \times 32 \times 3$). Each label or pixel is a 32-bit value in the memory. Refer to [CIFAR-10 website](#) for further details.

Working Items

1. Study the software version of the k-NN and complete it for performance comparison. Compute distance with the square of L2 distance (Euclidean distance) with $k=5$, i.e., $(L2)^2$. (We want to avoid the complicated square root operation.)
2. Implement the hardware version of k-NN using **either** PCPI or MMAP method.
3. You don't need to put everything into the hardware part. Instead, you can partition the algorithm into several steps (e.g. distance computation, sorting, ..., and so on). And let the software control when and which steps to execute.
4. Change the distance metric and the value of k . Try to achieve a better accuracy than the default setting. (We only consider a small subset of the images, so you may expect for a relatively low accuracy.)

Questions & Discussion

1. How do you partition your algorithm?
2. Profile both software and hardware versions. Which part do you think is the bottleneck of your (software/hardware) design?
3. Briefly describe the experiments you did and their results.